




# EDA(Exploratory Data Analysis)

## ✓ EDA(Exploratory Data Analysis)

EDA is essential for understanding the underlying patterns, relationships, and structures within data. It helps in detecting outliers, testing assumptions, and forming hypotheses. Without EDA, data analysis is like navigating in the dark.

1. Variable Identification
2. Univariate Analysis
3. Bivariate Analysis
4. Outlier Treatment
5. Missing Value Treatment
6. Variable Creation
7. Variable Transformation



```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df_raw=pd.read_excel("/content/drive/MyDrive/FSDS @Kodi Senapati/Datasets/Rawdata.xlsx")
```

```
df=pd.read_excel("/content/drive/MyDrive/FSDS @Kodi Senapati/Datasets/Rawdata.xlsx")
```

```
print(df)
```

```

  Name      Domain      Age  Location  Salary  Exp
0  Mike  Data science  34 years  Mumbai    50000    2+
1  Teddy  Testing    45' yr  Bangalore  100000    <3
2  Uma  Data analyst  NaN      NaN    150000  4+ yrs
3  Jane  Analytics  NaN      Hyderabad  20000    NaN
4  Uttam  Statistics  67-yr    NaN    30000-  5+ year
5  Kim    NLP        55yr      Delhi    60000$    10+

```

```
print(df.shape)
```

```
print(df.columns)
```

```

(6, 6)
Index(['Name', 'Domain', 'Age', 'Location', 'Salary', 'Exp'], dtype='object')

```

```
print(df.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6 entries, 0 to 5
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Name        6 non-null      object
1   Domain       6 non-null      object
2   Age          4 non-null      object
3   Location     4 non-null      object
4   Salary       6 non-null      object
5   Exp          5 non-null      object
dtypes: object(6)
memory usage: 416.0+ bytes
None

```

```
df.isnull().sum()
```

```

0
Name      0
Domain    0
Age       2
Location  2
Salary    0
Exp       1

```

Initially before applying EDA technique cleaning of data is required.

```
df['Name']
```

```

  Name
0  Mike
1  Teddy
2  Uma
3  Jane
4  Uttam
5  Kim

```

```
# Remove non word by replacing it with ''
```

```
df['Name']=df['Name'].str.replace(r'\W','',regex=True)
```

```
print(df['Name'])
```

```
0    Mike
1    Teddy
2    Umar
3    Jane
4    Uttam
5    Kim
Name: Name, dtype: object
```

```
df['Domain']
```

```
0    Datascience#$
1         Testing
2    Dataanalyst^^#
3    Ana^alytics
4         Statistics
5         NLP
```

```
df['Domain']=df['Domain'].str.replace(r'\W','',regex=True)
df['Domain']
```

```
0    Datascience
1         Testing
2    Dataanalyst
3     Analytics
4     Statistics
5         NLP
```

```
df
```

```
0    Mike  Datascience  34 years  Mumbai  5^00#0  2+
1    Teddy    Testing  45' yr  Bangalore  10%%000  <3
2    Umar  Dataanalyst   NaN      NaN  1$5%000  4> yrs
3    Jane    Analytics   NaN  Hyderabad  2000^0   NaN
4    Uttam  Statistics  67-yr      NaN  30000-  5+ year
5    Kim      NLP  55vr      Delhi  6000^$0  10+
```

Next steps:

[Generate code with df](#)
[View recommended plots](#)
[New interactive sheet](#)

```
df['Age']=df['Age'].str.replace(r'\W','',regex=True)
df['Age']
```

```
# This removes non word characters but years,yr are word characters
```



	Age
0	34years
1	45yr
2	NaN
3	NaN
4	67yr
5	55yr



```
df['Age']=df['Age'].str.extract('(\d+)')
df['Age']
```



	Age
0	34
1	45
2	NaN
3	NaN
4	67
5	55



```
df['Location']
```



	Location
0	Mumbai
1	Bangalore
2	NaN
3	Hyderbad
4	NaN
5	Delhi




```
df['Salary']
```



	Salary
0	5^00#0
1	10%%000
2	1\$5%000
3	2000^0
4	30000-
5	6000^\$0




```
df['Salary']=df['Salary'].str.replace(r'\W', "", regex=True)
df['Salary']
```



	Salary
0	5000
1	10000
2	15000
3	20000
4	30000
5	60000

```
df['Exp']
```




	Exp
0	2+
1	<3
2	4> yrs
3	NaN
4	5+ year
5	10+

```
df['Exp']=df['Exp'].str.replace(r'\W','',regex=True)
df['Exp']=df['Exp'].str.extract('(\d+)')
df['Exp']
```



	Exp
0	2
1	3
2	4
3	NaN
4	5
5	10

```
# Raw data
print("---Raw Data---")
print(df_raw)
print('\n')
print("---Clean Data---")
#Clean Data
print(df)
```



	Name	Domain	Age	Location	Salary	Exp
0	Mike	Datascience#\$	34 years	Mumbai	5^00#0	2+
1	Teddy^	Testing	45' yr	Bangalore	10%000	<3
2	Uma#r	Dataanalyst^^#	NaN	NaN	1\$5%000	4> yrs
3	Jane	Ana^alytics	NaN	Hyderbad	2000^0	NaN
4	Uttam*	Statistics	67-yr	NaN	30000-	5+ year
5	Kim	NLP	55yr	Delhi	6000^\$0	10+

	Name	Domain	Age	Location	Salary	Exp
0	Mike	Datascience	34	Mumbai	5000	2
1	Teddy	Testing	45	Bangalore	10000	3
2	Umar	Dataanalyst	NaN	NaN	15000	4
3	Jane	Analytics	NaN	Hyderbad	20000	NaN
4	Uttam	Statistics	67	NaN	30000	5
5	Kim	NLP	55	Delhi	60000	10

```
# Copy Clean data for applying EDA
```

```
cdf=df.copy()
print(cdf)
```

```
↗
```

	Name	Domain	Age	Location	Salary	Exp
0	Mike	Datascience	34	Mumbai	5000	2
1	Teddy	Testing	45	Bangalore	10000	3
2	Umar	Dataanalyst	NaN	NaN	15000	4
3	Jane	Analytics	NaN	Hyderbad	20000	NaN
4	Uttam	Statistics	67	NaN	30000	5
5	Kim	NLP	55	Delhi	60000	10

## ✓ 1. Variable Identification

It includes finding out the dependent variables(target variables) and independent variables(features variables). Finding out the relevant attributes and its types are to be considered to avoid multicollinearity, which may cause overfitting(causing less accuracy and high error).

```
print(cdf.columns)
```

```
↗ Index(['Name', 'Domain', 'Age', 'Location', 'Salary', 'Exp'], dtype='object')
```

```
# Identify dependent and independent variables
```

```
dep_var=cdf[['Salary']]
indep_var=cdf[['Name', 'Domain', 'Age', 'Location', 'Exp']]
```

```
print('Dependent Variable (Target)\n',dep_var)
print('\n')
print('Independent Variable (Features)\n',indep_var)
```

```
↗ Dependent Variable (Target)
Salary
0    5000
1   10000
2   15000
3   20000
4   30000
5   60000
```

```
Independent Variable (Features)
Name      Domain  Age  Location  Exp
0  Mike  Datascience  34    Mumbai    2
1  Teddy    Testing  45    Bangalore    3
2  Umar  Dataanalyst  NaN         NaN     4
3  Jane   Analytics  NaN    Hyderbad   NaN
4  Uttam  Statistics  67         NaN     5
5  Kim           NLP   55     Delhi    10
```

## ✓ 2. Univariate Analysis

It focuses on examining a single variable to understand its distribution, central tendency (e.g., mean, median), spread (e.g., variance, standard deviation), presence of outliers, and visualizations like histograms, box plots, and frequency tables. This analysis is crucial for gaining insights about each feature in the dataset individually.


### Univariate Analysis for Age

```
# Summary Statistics for Age
```


```
print("Summary statistics for Age:")
print(cdf['Age'].describe())
```

```
↗ Summary statistics for Age:
count      4
unique      4
top        34
freq        1
Name: Age, dtype: object
```

```
# Plotting histogram for Age
plt.figure(figsize=(5, 3))
sns.histplot(cdf['Age'], kde=True, bins=5)
plt.title('Distribution of Age')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```



```
# Plotting box plot for Age
plt.figure(figsize=(5, 3))
sns.boxplot(x=cdf['Age'])
plt.title('Box Plot of Age')
plt.xlabel('Age')
plt.show()
```




### Univariate Analysis for Salary

```
# Summary statistics for Salary
print("Summary statistics for Salary:")
print(cdf['Salary'].describe())
```




```
Summary statistics for Salary:
count      6
unique      6
top      5000
freq        1
Name: Salary, dtype: object
```

```
# Plotting histogram for Salary
plt.figure(figsize=(5, 3))
sns.histplot(cdf['Salary'], kde=True, bins=5)
plt.title('Distribution of Salary')
plt.xlabel('Salary')
plt.ylabel('Frequency')
plt.show()
```



```
# Plotting box plot for Salary
plt.figure(figsize=(5, 3))
sns.boxplot(x=cdf['Salary'])
plt.title('Box Plot of Salary')
plt.xlabel('Salary')
plt.show()
```



#### Explanation:

- `df['Age'].describe()` and `df['Salary'].describe()` provide summary statistics such as count, mean, standard deviation, minimum, quartiles, and maximum.
- `sns.histplot()` plots the histogram and overlays a KDE (kernel density estimate) to visualize the distribution.
- `sns.boxplot()` helps identify outliers and understand the spread and quartiles of the data.

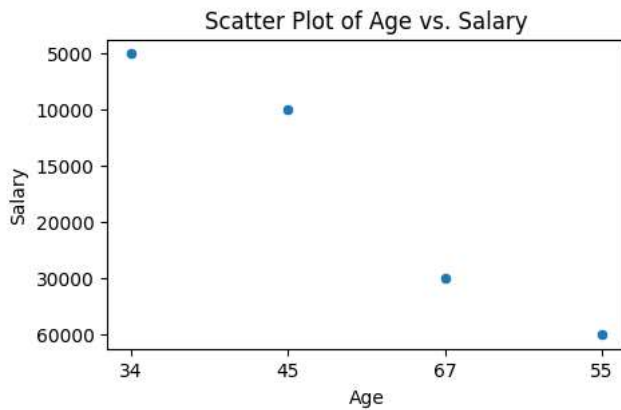
## ✓ 3. Bivariate Analysis

It is the analysis of two variables to explore the relationship between variables. For example, analyzing how Age relates to Salary or how Experience correlates with Salary.

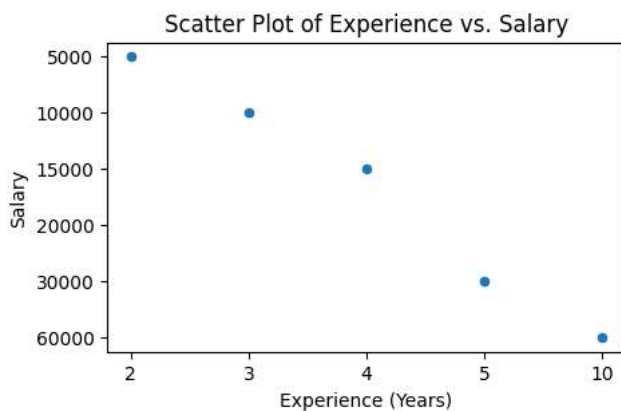
```
# Scatter plot for Age vs Salary
```

```
plt.figure(figsize=(5, 3))
sns.scatterplot(x='Age', y='Salary', data=cdf)
plt.title('Scatter Plot of Age vs. Salary')
plt.xlabel('Age')
plt.ylabel('Salary')
plt.show()
```





```
# Scatter Plot for Experience vs. Salary
plt.figure(figsize=(5, 3))
sns.scatterplot(x='Exp', y='Salary', data=cdf)
plt.title('Scatter Plot of Experience vs. Salary')
plt.xlabel('Experience (Years)')
plt.ylabel('Salary')
plt.show()
```



#### Explanation:

- `sns.scatterplot()` is used to visualize the relationship between two numerical variables (Age vs. Salary, Exp vs. Salary).
- This plot helps identify trends (e.g., if salary increases with age or experience).

```
# Correlation Matrix
correlation_matrix = cdf[['Age', 'Salary', 'Exp']].corr()
print("Correlation Matrix:")
print(correlation_matrix)
```




```
Correlation Matrix:
      Age  Salary  Exp
Age  1.000000  0.620110  0.552102
Salary 0.620110  1.000000  0.991064
Exp    0.552102  0.991064  1.000000
```

#### Explanation:

- `df.corr()` calculates the correlation between numerical variables.
- A value close to 1 indicates a strong positive correlation, while a value close to -1 indicates a strong negative correlation.
- A value around 0 indicates little or no linear correlation.

```
# Heatmap for correlation matrix
```

```
plt.figure(figsize=(4, 3))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', vmin=-1, vmax=1)
plt.title('Correlation Matrix Heatmap')
plt.show()
```

**Explanation:**

`sns.heatmap()` provides a visual representation of the correlation matrix, highlighting strong and weak relationships.

## ✓ 4. Outlier Treatment

It involves identifying and addressing extreme values in a dataset that deviate significantly from the other observations. Outliers can distort statistical analyses and lead to inaccurate insights or model performance. Common methods for treating outliers include:

- Removing Outliers: Excluding data points that fall outside a certain range (e.g., beyond 1.5 times the interquartile range (IQR)).
- Capping Outliers: Setting a maximum or minimum value to limit the impact of outliers.
- Transformation: Applying transformations like log or square root to reduce the impact of outliers.

**Treat outliers using the Interquartile Range (IQR) method, with an example of handling noisy data.**

**What is the IQR Method?**

The Interquartile Range (IQR) is a measure of statistical dispersion and is used to identify outliers in a dataset. It is calculated as the difference between the third quartile (Q3) and the first quartile (Q1):

$$\text{IQR} = Q3 - Q1$$

Outlier bounds:

$$\text{Lower Bound: } Q1 - 1.5 * \text{IQR}$$

$$\text{Upper Bound: } Q3 + 1.5 * \text{IQR}$$

Any data point outside these bounds is considered an outlier.

cdf

	Name	Domain	Age	Location	Salary	Exp
0	Mike	Datascience	34	Mumbai	5000	2
1	Teddy	Testing	45	Bangalore	10000	3
2	Umar	Dataanalyst	NaN	NaN	15000	4
3	Jane	Analytics	NaN	Hyderbad	20000	NaN
4	Uttam	Statistics	67	NaN	30000	5
5	Kim	NLP	55	Delhi	60000	10

Next steps:

[Generate code with cdf](#)

[View recommended plots](#)

[New interactive sheet](#)

# Function to treat outliers using IQR

```
def treat_outliers_iqr(column):
    Q1 = cdf[column].quantile(0.25)
    Q3 = cdf[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
```

# Replace outliers with boundary values

```
# Replace outliers with boundary values
cdf[column] = cdf[column].apply(lambda x: upper_bound if x > upper_bound else (lower_bound if x < lower_bound else x))

# Treat outliers in the numerical columns
for col in ['Age', 'Salary', 'Exp']:
    cdf[col] = pd.to_numeric(cdf[col], errors='coerce') # Ensure numeric dtype for calculation
    treat_outliers_iqr(col)

print("Updated DataFrame:")
print(cdf)
```

```
Updated DataFrame:
   Name  Domain  Age  Location  Salary  Exp
0  Mike  Datascience  34.0    Mumbai  5000.0  2.0
1  Teddy   Testing  45.0    Bangalore  10000.0  3.0
2  Umar  Dataanalyst   NaN         NaN  15000.0  4.0
3  Jane   Analytics   NaN    Hyderabad  20000.0  NaN
4  Uttam  Statistics  67.0         NaN  30000.0  5.0
5  Kim      NLP      55.0      Delhi  51875.0  8.0
```

## 5. Missing Value Treatment

It is essential in data preprocessing to ensure the dataset is complete and suitable for analysis or modeling. Missing values can lead to biased results and affect model performance. The approach to handling missing values depends on the nature of the data and the percentage of missing values.

### Common Strategies for Treating Missing Values:

#### 1. Deletion

- Listwise Deletion: Remove rows with missing values (use when missing data is minimal and does not impact analysis).
- Column Deletion: Remove columns with a large proportion of missing values.

#### 2. Imputation

- Mean/Median/Mode Imputation: Replace missing values with the mean, median, or mode (useful for numerical data).
- Forward/Backward Fill: Fill missing values based on neighboring data.
- Custom Imputation: Use a specific value or domain knowledge for imputation.

#### 3. Advanced Techniques

- Predictive Imputation: Use machine learning models to predict and fill missing values.
- K-Nearest Neighbors (KNN): Impute based on similar data points.

```
# Find the null values
```

```
print(cdf.isnull().sum())
```

```
Name      0
Domain     0
Age        2
Location   2
Salary     0
Exp        1
dtype: int64
```

```
cdf[['Age', 'Location', 'Exp']].info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6 entries, 0 to 5
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Age         4 non-null      float64
1   Location    4 non-null      object
2   Exp         5 non-null      float64
dtypes: float64(2), object(1)
memory usage: 272.0+ bytes
```

```
# Convert 'Age' and 'Exp' columns to numeric, coercing errors to NaN
```

```
cdf['Age'] = pd.to_numeric(cdf['Age'], errors='coerce')
```

```
cdf['Exp'] = pd.to_numeric(cdf['Exp'], errors='coerce')
```

```
#Impute the missing values with the median
cdf['Age'].fillna(cdf['Age'].median(),inplace=True)
cdf['Exp'].fillna(cdf['Exp'].median(),inplace=True)
```

 <ipython-input-202-694922a1ee28>:7: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assign. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting value is a copy.


For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value, inplace=True)

```
cdf['Age'].fillna(cdf['Age'].median(),inplace=True)
<ipython-input-202-694922a1ee28>:8: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assign. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting value is a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value, inplace=True)

```
cdf['Exp'].fillna(cdf['Exp'].median(),inplace=True)
```

```
cdf.isnull().sum()
```




	0
Name	0
Domain	0
Age	0
Location	2
Salary	0
Exp	0

dtype: int64

```
# Fill the Location blank values
```

```
# Impute 'Location' using mode
cdf['Location'].fillna(cdf['Location'].mode()[0], inplace=True)
```

 <ipython-input-204-48775de0713e>:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assign. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting value is a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value, inplace=True)

```
cdf['Location'].fillna(cdf['Location'].mode()[0], inplace=True)
```

```
cdf.isnull().sum()
```



	0
Name	0
Domain	0
Age	0
Location	0
Salary	0
Exp	0

dtype: int64

```
# Save the clean data as a separate sheet
```

```
cdf.to_csv('clean_data.csv')
```

```
print(cdf)
```

```

➡
  Name      Domain  Age  Location  Salary  Exp
0  Mike  Datascience  34.0  Mumbai    5000.0  2.0
1  Teddy    Testing  45.0  Bangalore 10000.0  3.0
2  Umar  Dataanalyst  50.0  Bangalore 15000.0  4.0
3  Jane    Analytics  50.0  Hyderabad 20000.0  4.0
4  Uttam  Statistics  67.0  Bangalore 30000.0  5.0
5  Kim      NLP      55.0  Delhi    51875.0  8.0

```

## 6. Variable Creation

This involves generating new columns based on existing data, which can provide additional information or create more meaningful categorizations.

```

# Apply one-hot encoding to the 'Domain' and 'Location' columns
var_cre_df = pd.get_dummies(cdf, columns=['Domain', 'Location'], prefix=['Domain', 'Loc'])

```

```

print("Updated DataFrame with One-Hot Encoded Variables:")
print(var_cre_df)

```

```

➡ Updated DataFrame with One-Hot Encoded Variables:
   Name  Age  Salary  Exp  Domain_Analytics  Domain_Dataanalyst  \
0  Mike  34.0  5000.0  2.0             False                False
1  Teddy  45.0 10000.0  3.0             False                False
2  Umar   50.0 15000.0  4.0             False                 True
3  Jane   50.0 20000.0  4.0              True                False
4  Uttam  67.0 30000.0  5.0             False                False
5  Kim   55.0 51875.0  8.0             False                False

   Domain_Datascience  Domain_NLP  Domain_Statistics  Domain_Testing  \
0                True          False             False             False
1                False          False             False              True
2                False          False             False             False
3                False          False             False             False
4                False          False              True             False
5                False           True             False             False

   Loc_Bangalore  Loc_Delhi  Loc_Hyderabad  Loc_Mumbai
0             False       False           False          True
1              True       False           False          False
2              True       False           False          False
3             False       False            True          False
4              True       False           False          False
5             False       True            False          False

```

```
print(var_cre_df.columns)
```

```

➡ Index(['Name', 'Age', 'Salary', 'Exp', 'Domain_Analytics',
        'Domain_Dataanalyst', 'Domain_Datascience', 'Domain_NLP',
        'Domain_Statistics', 'Domain_Testing', 'Loc_Bangalore', 'Loc_Delhi',
        'Loc_Hyderabad', 'Loc_Mumbai'],
        dtype='object')

```

## 7. Variable Transformation

Label Encoding is typically used for converting categorical variables into numerical representations, where each unique category is assigned an integer value. While Label Encoding is mostly used for categorical data, it can also be considered a form of variable transformation when transforming text or categorical data into a numerical format that can be used by machine learning algorithms.

```
from sklearn.preprocessing import LabelEncoder
```

```

# Initialize the LabelEncoder
label_encoder = LabelEncoder()

```

```
Var_trans_df = cdf.copy()
```

```

# Apply Label Encoding to 'Domain' column
Var_trans_df['Domain_Encoded'] = label_encoder.fit_transform(Var_trans_df['Domain'])

```

```
print(Var_trans_df)
```

```

➡
  Name      Domain  Age  Location  Salary  Exp  Domain_Encoded
0  Mike  Datascience  34.0  Mumbai    5000.0  2.0             2

```

1	Teddy	Testing	45.0	Bangalore	10000.0	3.0	5
2	Umar	Dataanalyst	50.0	Bangalore	15000.0	4.0	1
3	Jane	Analytics	50.0	Hyderbad	20000.0	4.0	0
4	Uttam	Statistics	67.0	Bangalore	30000.0	5.0	4
5	Kim	NLP	55.0	Delhi	51875.0	8.0	3

Start coding or [generate](#) with AI.